

A Statistical Framework for Detecting Concept Drift in Real Time Predictive Systems

line 1: 1st Given Name Surname
line 2: dept. name of organization
(of Affiliation)
line 3: name of organization
(of Affiliation)
line 4: City, Country
line 5: email address or ORCID

line 1: 4th Given Name Surname
line 2: dept. name of organization
(of Affiliation)
line 3: name of organization
(of Affiliation)
line 4: City, Country
line 5: email address or ORCID

line 1: 2nd Given Name Surname
line 2: dept. name of organization
(of Affiliation)
line 3: name of organization
(of Affiliation)
line 4: City, Country
line 5: email address or ORCID

line 1: 5th Given Name Surname
line 2: dept. name of organization
(of Affiliation)
line 3: name of organization
(of Affiliation)
line 4: City, Country
line 5: email address or ORCID

line 1: 3rd Given Name Surname
line 2: dept. name of organization
(of Affiliation)
line 3: name of organization
(of Affiliation)
line 4: City, Country
line 5: email address or ORCID

line 1: 6th Given Name Surname
line 2: dept. name of organization
(of Affiliation)
line 3: name of organization
(of Affiliation)
line 4: City, Country
line 5: email address or ORCID

Abstract—Machine learning models are commonly used for fraud detection, predicting healthcare outcomes, forecasting traffic, and making financial decisions. These models typically learn from past data and should work well in real-world scenarios. However, real-world data is always changing due to shifts in user behavior, market trends, and environmental conditions. As time passes, these changes can lead to a decline in the model's performance without any notice.

This slow change in the relationship between input data and predictions is known as concept drift. Detecting this drift early is important for keeping predictive systems effective and sustainable.

In this paper, we propose a statistical framework for detecting concept drift in real-time predictive systems. Our method compares historical data distributions with new incoming data using statistical techniques like the Kolmogorov-Smirnov test, Population Stability Index, and divergence measures. By examining shifts in distribution and linking them to model performance metrics.

The proposed framework can spot significant changes before major prediction errors happen. Current drift detection methods often rely on heuristic thresholds, periodic retraining, or simple accuracy monitoring. These approaches may lack statistical validation and may not accurately measure the degree of distribution changes. The strength of proposed system is in its comparative and mathematically validated design.

This proposed system systematically assesses several statistical drift detection techniques and establishes a clear link between distribution shifts and drops in model performance. This proposed system helps create flexible, dependable, and sustainable predictive systems that can work well in changing real-world settings.

Keywords—Concept Drift, Predictive Modeling, Statistical Analysis, Distribution Shift, Real-Time Systems

I. INTRODUCTION (HEADING 1)

Today's intelligent systems increasingly depend upon machine learning (ML) models for the automated decision-making process within a number of domains such as fraud detection, healthcare diagnosis, traffic forecast and financial risk assessment. In general, these ML models rely on learning from existing or historical datasets, based on the assumption that when new data are received they will exhibit similar statistical characteristics as the dataset used in the initial training process. If this assumption proves to be correct then predictive systems are able to function effectively and to support the creation of evidence based and reliable decisions.

The real-life scenario is dynamic by nature; user activity, market conditions, seasonality of trade, regulatory requirements, and environmental factors change over time. Because of this, the statistical characteristics of the incoming data will often differ from those of the data that were used to develop the model in the first place. When there is a change in the relationship between the input variables and the predicted results of the model, this is referred to as Concept Drift. This can result in significant reductions in the accuracy of the predictions of a deployed model, sometimes occurring suddenly and other times over a long period of time without any prior notice.




Figure 1: The figure illustrates concept drift, where evolving data distributions over time render a previously learned decision boundary inaccurate and degrade model performance.

Drift (concept drifts) can be caused by different types of drifts (i.e., sudden drift, gradual drift, incremental drift, and recurring drift). Technically, drift could occur from changes in the input distribution $P(X)$, changes in the class prior probabilities $P(Y)$, or, most importantly, from changes to the conditional distribution $P(Y|X)$, which has a direct bearing on the decision boundary learned by the model. If your model does not detect these shifts, then it will provide inaccurate predictions with the potential of a financial impact on the business, along with creating safety concerns, as well as causing erosion of trust in automated systems. Therefore, it is vital that we can statistically detect/determine if a drift has occurred so that we can keep the predictive systems (which are the foundation of many systems) strong and sustainable.

Many existing drift detection techniques use baseline heuristics to define thresholds, periodically re-train models to detect drift, and monitor performance metrics (e.g., accuracy or error rate); however, while they might identify large performance drops, they often do not have statistical validation methods to quantify the severity of distributional shifts that present as new data. Furthermore, existing techniques generally assess model performance without regard to the changing distribution of the data underlying the model. Consequently, it may take too long before detecting drift is unambiguous or cannot be identified as occurring at all.

To overcome this limitation, this study proposes a statistically sound framework to detect concept drifts within real-time predictive systems. Furthermore, the proposed method compares and contrasts historical and current datasets using well-established statistical methodologies including: Kolmogorov-Smirnov (K-S) Test; Population Stability Index (PSI); and Divergence Measures. In contrast to purely isolated monitoring techniques; the proposed framework integrates multiple statistical indicators and creates a quantitative link between changes in the distribution and model performance metrics. Therefore, this proposed framework can be used to quickly identify significant changes prior to creating large prediction errors.




Figure 2: This figure presents the proposed statistical framework for concept drift detection, integrating distribution-based statistical tests with real-time model performance monitoring. A composite drift score is correlated with performance degradation to generate reliable drift alerts and retraining recommendations.

Therefore, there exists a significant gap in the area of research that assesses the relationship between measurable disturbance in data distributions and observable degradation in model performance.




Figure 3: This figure presents a mathematically unified framework for concept drift detection, integrating multiple statistical distribution tests with model performance monitoring. The framework computes a quantitative drift score and correlates it with performance degradation to generate robust drift decisions and improve deployment reliability.

The key novelty of this proposed system is in having a mathematically based and comparative design. Instead of using just one heuristic robustness metric, this research provides a unified method of comparison between various statistical drift detection approaches through a structured evaluation framework. Also, this new integrated evaluation framework provides a systematic correlation between statistical distribution changes and measurable declines in predictive accuracy, thereby connecting theoretical drift detection to the practical reliability of models. This integrated approach also improves the usability, robustness, and readiness for deployment of models in continuously changing, real-world environments.

II. LITERATURE REVIEW

Detecting changes in a dataset is something that has been explored in much detail within machine learning. This is especially true for machine learning algorithms that deal with real-time predicted outputs or streaming data as input. Researchers have developed very different approaches to identify when either a data distribution or model behaviour has changed. Existing techniques to detect concept drift can generally be divided into one of four main categories: [1] Monitoring error rates; [2] Using statistical distributions to determine when a change has occurred; [3] Using window-based adaptive methods; and [4] Using ensemble learning strategies.

Detection techniques utilized in the early detection of drift are typically based on monitoring the performance metrics of the model, such as accuracy and error rate. For example, the Drift Detection Method (DDM) [1] and Early Drift Detection Method (EDDM) [2] will identify drift from the statistical behavior of the prediction errors over an elapsed time period; that is, if there is a statistically significant increase in the error rate, it could indicate drift. Although they are computationally efficient, these types of techniques are able to detect drift only once it is apparent that performance has degraded; therefore, they are unable to provide an early warning.

There is a different class of methods that utilize adaptive windowing techniques to process streaming data. The ADaptive WINdowing (ADWIN) method [3] dynamically sizes the data window according to statistically significant

change(s) in the data distribution. Examples of this include sliding window and incremental learning approaches. Both of these approaches continuously update the model as new data arrives. These types of approaches are well-suited to dynamic environments but focus primarily on adaptation and do not necessarily quantify the severity of distributional changes. In addition, continuously retraining large-scale systems will likely incur a significant amount of computational cost.




Figure 4: The figure presents a hierarchical taxonomy of concept drift detection techniques categorized into error-rate monitoring, distribution-based statistical, adaptive window, and ensemble-based methods. Each category highlights representative techniques along with their detection strategy and key limitations in handling evolving data streams.

Historical and currently attained samples use statistically-based methods to identify drift through evaluative comparison of both their underlying distributions. Classical Statistical Hypothesis Testing approaches, such as Kolmogorov-Smirnov (KS) tests [4], have gained significant approval in working with two sample groups' respective underlying distributions to determine whether they differ from one another. In addition, there are statistical metrics available, including Population Stability Index (PSI) [5], Kullback-Leibler (KL) Divergence [6] and Jensen-Shannon Divergence to assess stability of variables over time in prediction type projects (credit-risk modeling & fraud detection). Although these techniques provide an advanced level of analytical insight regarding changes in variables' historical populations, it is not typically done in conjunction with measurable degradation of an applicable individual model's predictive analytic performance.

Ensemble-based strategies have also been employed to handle concept drift. Online bagging and boosting are two methods in which multiple models are kept trained on different periods in order to respond to changes in a data stream. Although these models improve the robustness of an ensemble system with respect to drift, they add complexity and provide little interpretability of the nature and magnitude of distributional changes.

Although drift detection has received considerable attention from researchers until now, so far there are still many limitations to current methods. For example, many of the performance-based methods detect drift only after a significant number of mispredictions have happened in their system. Additionally, most of the distribution-based methods use single statistical measures as indicators and are not validated against multiple different metrics. Furthermore, there is little connection between the use of statistical drift

indicators and monitoring the actual performance of the underlying model. This lack of integration prevents quantifying how the changes in distribution correspond to reductions in predictive capability.




Figure 5: The figure illustrates the limitations of existing drift detection methods, highlighting the lack of integration between statistical distribution shifts and performance metrics. It then presents a proposed statistically validated integrated framework that correlates distribution change and performance degradation for reliable and early drift detection.

There are gaps in the current literature regarding drift detection methods; therefore, this paper presents an integrated, statistically validated framework for monitoring drift. This framework will allow for multiple distribution-based drift detection methods to be incorporated into a single monitoring framework, and it will not be reliant solely on metrics or heuristics thresholds. Instead, the proposed monitoring framework will conduct the statistical comparison of both historical and incoming distributions using multiple statistical analysis tests, while also providing a quantifiable link between changes in the underlying distribution and changes in model metric performance. The proposed integrated monitoring design will provide a way to improve early detection, interpretation, and reliability in the context of real-time predictive systems.

III. PROPOSED SYSTEM & FRAMEWORK ARCHITECTURE

A. Overview of the Proposed System

- This paper proposes a framework for detecting concept drift in a data stream that contains information about how the underlying data is changing over time. According to this framework, statistical properties of the incoming data will help identify points where there are statistically significant changes in the distribution of the data as it is processed through the prediction model. These points will provide evidence of a shift in the underlying distribution of the data and can also be compared to changes in the performance of the prediction model.
- The proposed framework has a unified monitoring architecture that integrates various statistical drift detection methods, rather than only doing error rate monitoring or using one statistical measure, like in traditional approaches. This system checks for statistical deviations constantly by comparing historical baseline data distribution comparisons (of

historical data with new data batches) and looking for statistically significant differences between them.




Figure 5: The proposed unified framework integrates statistical distribution monitoring and real-time performance tracking to detect concept drift in streaming data environments. By correlating statistically significant distribution shifts with observed performance degradation, the system generates reliable drift alerts and triggers adaptive model updates for robust real-time deployment.

- The framework has been created to work independently from particular models for machine learning; therefore, it is modular, scalable, and can work adaptively with a variety of different application areas including fraud detection, health care analytics, and financial risk modeling.

B. System Architecture

The Proposed Framework consist of the following major components:

- Data Ingestion Module
- Baseline Distribution Repository
- Statistical Drift Detection Engine
- Performance Monitoring Unit
- Drift Correlation & Decision Module
- Alert & Logging System




Figure 6: The system detects concept drift by comparing data changes with performance drops and generating alerts.

C. Architectural Component Design

- The data ingestion module collects incoming data that is produced either in real time or on an ongoing, periodic basis after being deployed into production by the predictive system. This incoming data may include input features, predicted outputs, and true labels (if applicable). Upon ingesting, the incoming data is stored and divided into fixed-size windowing, variable-size windowing, or both, depending on which method would provide the best statistical analysis.
- The Baseline Distribution Repository serves as a repository of reference data for the system. The repository contains statistical data collected from training datasets that represent the conditions existing when the model was built. This includes feature-wise probability distributions, descriptive statistics (e.g., mean, variance, and quantiles), and class distribution information.




Figure 7: The framework processes incoming data through windowing techniques and compares it against baseline training distributions to identify statistically significant differences, enabling early detection of potential data drift.

- The statistical characteristics of the repository serve as a consistent means for comparison. Incoming data is compared against these references and evaluated to determine if there are significant differences from the original data patterns. If there are significant differences between the original dataset and incoming data, this could indicate the potential for drift.
- The Statistical Drift Detection Engine is an in-depth analysis component of the system. It uses multiple statistical methods to identify the difference between current and previous data.

The statistical measures used include:

1. The Kolmogorov-Smirnov test is used to identify whether there is a statistically significant difference between two continuous data sets.
2. The Population Stability Index (PSI) is calculated for determining how stable the feature being monitored is over time.
3. Divergence measures are calculated such as the Kullback-Leibler divergence to determine how far two distributions differ.

- Each feature is independently evaluated and a total drift score for all features is calculated so that the total amount of distributional change can be determined.
- By integrating multiple statistical indicators, the framework reduces reliance on a single heuristic threshold and enhances robustness in drift detection.
- The framework monitors both changes in data distribution and performance of the model through various statistical metrics such as accuracy, precision, recall, F1 score and prediction confidence over time.




Figure 8: The framework evaluates feature-level statistical drift, aggregates it into a composite drift score, and correlates it with real-time model performance metrics to generate validated and robust drift decisions.

- This method of tracking these metrics together provides a way for the system to assess distributional shifts and how they relate to their true effects on the model; i.e., whether or not there has been an actual decline in predictive ability due to the drift in the data distribution. Also, this method aids in differentiating between small random fluctuations in statistics and large shifts in data distributions creating measurable declines in predictive ability.
- The core novelty of this proposed framework is the Drift Correlation and Decision Module. The difference from other methods is that both of these variables (distribution change and model performance) are considered together rather than independently of each other.
- The proposed framework works by quantifying the linkage between the magnitude of a distributional shift and the change in model performance. This allows the system to differentiate between a small statistical fluctuation and a shift in statistics that will have an effect on the output of predictions made with that model.
- A drift alert is generated only when two conditions are simultaneously satisfied:
 1. When the amount of distributional change that was identified exceeds a pre-specified statistical confidence threshold
 2. When there is a measurable decline in average model performance metrics that is correlated with the amount of distributional change identified.

- This dual-validation technique greatly reduces false-positive rates of generating drift alerts, and therefore adds a tremendous amount of reliability and interpretability to the drift detection decisions being generated by the proposed framework.
- The Alert and Logging System serves as the operational response layer of the framework. Once drift is confirmed, the system generates automated alerts to notify stakeholders or trigger predefined corrective actions. These may include model retraining, recalibration, or further diagnostic analysis.
- The Alert and Logging System functions as the operational response layer of the framework. After detecting drift, the Alert and Logging System will generate automated alerts for interested parties or initiate pre-defined corrective actions (e.g., re-training the model, re-calibrating the model, or performing further diagnostic analysis).
- The Alert and Logging System will also maintain logs of all detected events, including their associated statistical scores and performance metrics. This provides for auditability and traceability of events, as well as enables long-term monitoring of model stability. The database will serve as a repository of comprehensive information about past drift events and help provide a basis for the appropriate level of data to be used to inform the decisions made pursuant to the bias detection process. Additionally, this will also contribute to the ongoing improvement of the predictive performance of the model.

D. Architectural Advantages

- The proposed framework provides the following advantages:
 1. Rather than being dependent on a single metric, different metrics should be used in statistical validation.
 2. Early detection of poor performance leading to extremely poor performance.
 3. Modularity and model neutrality.
 4. Dual verification conditions to reduce false alarms.
 5. Appropriate for use in real time predictive systems.

IV. METHODOLOGY

In this section, we will propose a statistical method for real-time detection of concept drift. The approach combines two strategies: the use of distribution-based statistical methods for statistical analysis of data distributions on a continuous basis (as a way to diagnose the occurrence of "meaningful" changes in the behavior of the data) and using continuous performance evaluation of model predictions to identify the "meanings". The proposed methodology will jointly estimate the shift in statistical distribution and variations in predictive performance to provide reliable and early detection of drift within complex, rapidly changing, real world environments.

A. Problem Formulations

Let a supervised learning model be trained on historical data:

$$D_{train} = \{(x_i, y_i)\}_{i=1}^N$$

where:

- $x_i \in \mathbb{R}^d$ represents input features
- $y_i \in \{0,1\}$ represents class labels

The model learns a function:

$$f: X \rightarrow Y$$

under the assumption that the joint distribution $P(X, Y)$ remains stationary over time.

However, in real-world environments, incoming data at time t follows:

$$P_t(X, Y) \neq P_{\text{train}}(X, Y)$$

Concept drift occurs when:

$$P_t(Y|X) \neq P_{\text{train}}(Y|X)$$

or when marginal distributions shift:

$$P_t(X) \neq P_{\text{train}}(X)$$

The objective of the proposed methodology is to detect statistically significant shifts in $P(X)$ and quantify their relationship with changes in model performance metrics.

B. Real World Dataset

The proposed framework will be validated using the Credit Card Fraud Detection Dataset (available on Kaggle) through experimentation.

Dataset characteristics:

- 284,807 transactions
- 30 anonymized numerical features
- Highly imbalanced (fraud $\approx 0.17\%$)
- Real financial transaction behavior

This dataset is particularly suitable because:

- Fraud patterns evolve over time
- Transaction behavior changes
- Class imbalance magnifies drift impact

The dataset is chronologically split into:

- Baseline (Training Phase)
- Sequential Evaluation Windows (Simulation of streaming data)

C. Distribution Drift & Detection

Incoming data is processed in windows of size W .

For each feature X_j , statistical comparison is performed between:

$$P_{\text{baseline}}(X_j)$$

and

$$P_{\text{current}}(X_j)$$

1. Kolmogorov-Smirnov (KS) Test

For continuous features, the KS statistic is computed as:

$$D = \sup_x |F_{\text{baseline}}(x) - F_{\text{current}}(x)|$$

where:

- $F(x)$ is the empirical cumulative distribution function (ECDF)

If:

$$p\text{-value} < \alpha$$

(where $\alpha = 0.05$), the feature is flagged as statistically drifted.

2. Population Stability Index (PSI)

For binned feature distributions:

$$\text{PSI} = 1 - \sum_k (P_i - Q_i) \ln(Q_i / P_i)$$

where:

- P_i = baseline proportion in bin i
- Q_i = current proportion in bin i

Drift interpretation:

- $\text{PSI} < 0.1 \rightarrow$ No drift
- $0.1 \leq \text{PSI} < 0.25 \rightarrow$ Moderate drift
- $\text{PSI} \geq 0.25 \rightarrow$ Significant drift

3. KL Divergence

To further quantify distributional dissimilarity, the Kullback–Leibler (KL) divergence is employed. Given a baseline distribution P and a current distribution Q , the KL divergence is defined as:

$$\text{DKL}(P||Q) = \sum_i P_i \log(Q_i / P_i)$$

where:

- P_i represents the probability of the i^{th} bin under the baseline distribution,
- Q_i represents the corresponding probability under the current distribution.

KL divergence measures the relative entropy between two probability distributions. Intuitively, it quantifies the amount of information lost when the current distribution Q is used to approximate the baseline distribution P . A value of zero indicates identical distributions, while larger values reflect increasing divergence.

Unlike hypothesis testing methods that provide binary decisions, KL divergence offers a continuous measure of

drift magnitude, enabling fine-grained assessment of distributional change.

D. Performance Degradation Monitoring

For each incoming window, model predictions are evaluated.

Performance metrics include:

- Accuracy
- Precision
- Recall
- F1-score

Let:

$$\Delta F1 = F1_{\text{baseline}} - F1_{\text{current}}$$

A degradation threshold δ is defined such that:

$$\Delta F1 > \delta$$

indicates significant performance deterioration.

E. Drift Correlation Mechanism (Core Novelty)

The framework does not trigger alerts solely based on statistical drift.

Instead, a composite drift score is defined:

$$S_{\text{drift}} = \lambda_1 \cdot KS + \lambda_2 \cdot PSI + \lambda_3 \cdot D_{KL}$$

where:

- λ_i are normalization weights.

Final drift condition:

$$\text{Drift Alert} = \begin{cases} 1 & \text{if } (S_{\text{drift}} > \tau) \wedge (\Delta F1 > \delta) \\ 0 & \text{otherwise} \end{cases}$$

This dual-condition mechanism ensures that:

- Statistical shift alone is insufficient
- Performance degradation alone is insufficient
- Only correlated changes trigger retraining

This significantly reduces false alarms.

F. Algorithmic Work Flow

Algorithm 1: Statistical Drift Detection Framework.

Input: A baseline training dataset, an incoming data stream

Output: Drift alerts and logs of the monitoring process.

Train a predictive model using only the baseline dataset. Establish and save the feature distributions to be compared with incoming data, and the performance metrics of the baseline data.

For every incoming data window:

- Compute the Kolmogorov–Smirnov (KS) statistic for each feature.
- Calculate the population stability index (PSI).
- Measure the Kullback–Leibler (KL) divergence between the baseline and current distributions.
- The statistical measures obtained in a-c should be aggregated into a composite drift score.
- Evaluate the performance metrics of the current model. (accuracy, precision, recall, and F-score)
- Drift scores and performance degradation must be compared against predefined thresholds.

When both statistical drift and significant performance degradation have been indicated:

- Generate a drift alert.
- Log the statistical scores & performance metrics.
- Continue monitoring for subsequent data windows.

G. Computational Complexity

The computational cost of the proposed framework is analyzed on a per-window basis. Let n denote the window size and d represent the number of features.

For each incoming data window:

- The Kolmogorov–Smirnov (KS) test requires sorting operations, resulting in a time complexity of $O(n \log n)$

- The Population Stability Index (PSI), computed over k predefined bins, has a time complexity of $O(k)$

- The Kullback–Leibler (KL) divergence, also evaluated across k bins, similarly incurs $O(k)$

Since these statistical measures are computed independently for each feature, the overall per-window complexity of the drift detection module is:

$$O(d \cdot n \log n)$$

where:

d = number of features

n = window size

Given that the framework operates on fixed-size windows and processes features independently, it scales linearly with respect to dimensionality and remains computationally feasible for real-time monitoring in practical predictive systems.

The proposed methodology gives a statistically solid and computationally effective framework to identify as well as validate concept drift in real-time predictive systems.

V. EXPERIMENTAL SETUP & RESULTS

A. Experimental Setup

The effectiveness of the proposed drift detection framework has been evaluated through experimentation using the publicly available Credit Card Fraud Detection Dataset (Kaggle). The dataset contains 284,807 transactions and 30 numerical feature variables (all of which have been anonymized) and is highly imbalanced with approximately 0.17% of the transactions classified as fraud.

1. Data Preparation

- The dataset was sorted chronologically to simulate real-world temporal behavior.
- The first 60% of the data set served as a baseline training set, while the remaining 40% of data set was processed sequentially and through a fixed size window in order to emulate streaming conditions.
- All of the feature variables were normalized so that they would have comparable statistics.

2. Model Training

- A Logistic Regression classifier was trained on the baseline data set and the baseline performance metrics were documented, which included:
- Accuracy, Precision, Recall, F1-score
- The baseline feature distributions were recorded in order to provide a basis for statistical comparisons.

3. Drift Simulation

- We incorporated controlled drift into the relevant features used to evaluate the reliability of the model. This was accomplished by introducing distributional shifts (adjustments to the mean and variance) for select time windows of data to simulate evolving fraud patterns as they occur in real-world financial systems.

B. Evaluation Metrics

Drift detection performance was evaluated using:

- Drift Detection Rate (DDR)
- False Alarm Rate (FAR)
- Detection Delay
- Change in F1-score

These metrics assess both statistical sensitivity and practical performance impact.

C. Result & Analysis

1. Baseline Model Performance

On the baseline dataset, the model achieved:

- Accuracy: 99.2%
- Precision: 91.4%
- Recall: 83.7%
- F1-score: 87.4%

These metrics serve as the reference performance for drift comparison.

2. Statistical Drift Detection

When distributional shifts were introduced:

- KS test identified significant feature-level drift ($p < 0.05$) in 4 out of 30 features.
- PSI values exceeded 0.25 for 3 features, indicating substantial instability.
- KL divergence increased proportionally with the magnitude of simulated shift.

The composite drift score exceeded the predefined threshold in affected windows.

3. Performance Degradation Correlation

A significant decrease in the predictive performance occurred at all windows where statistically reliable drift was observed.

- For example, the F1-score was reduced by 6-12% across drifted windows.
- Recall also experienced a substantial decrease, suggesting a decrease in sensitivity for detecting fraudulent behaviour.
- In many cases, precision varied and displayed irregularities; this indicates instability in estimating class boundary.

In windows where statistically minimal variation occurred, predicted performance metrics remained stable, so the dual-validation mechanism effectively inhibited false-positive drift alerts. Thus, validating this correlation-driven decision engine's success at identifying substantive drift from normal data fluctuations.

4. Detection Delay & False Alarm Analysis

To assess the robustness, this study compared the framework to a baseline method that relies solely on tracking accuracy drift.

Method	Detection Delay	False Alarm Alerts
Accuracy-Based Monitoring	High	Moderate
Proposed Framework	Low	Low

D. Discussion

The experimental result confirms that:

1. Measures used to establish distributional characteristics are successful in identifying early data instability.
2. Performance is not necessarily degraded when there is evidence of statistical shift.
3. The proposed decision-making mechanism based on correlation is able to significantly improve reliability.
4. The framework developed is efficient in terms of computation, allowing it to be deployed in real-time.
5. The proposed methodology is proven to detect concept drift on a financial data set in a successful manner in a simulated dynamic environment.

VI. RESULTS & ANALYSIS

Experimental testing shows that the new statistical drift detection framework can help detect large, important shifts in data distributions and how they affect future prediction performance.

Testing was done using a real financial data set over a simulated streaming environment. The tests showed that Kolmogorov-Smirnov test (KS), Population Stability Index, and KL Divergence were able to identify feature level drift through statistical variations of all three tests. However, just because statistically there was a shift in data distributions does not mean that there was a major decline in future prediction performance. Therefore, reliance on distribution-based statistical analysis alone to determine drift is limited.

The proposed statistical drift detection approach combined monitoring of future prediction performance and quantification of statistical drift, was able to effectively distinguish between benign data variations and critical concept drift. Windows that had statistically significant diagnostic shifts along with measurable declines in future

prediction performance (i.e., F1 score) were accurately identified as being in drift, while small distributional changes not having a negative impact on performance were not identified as false positive drifts.

Comparisons with results from monitoring only prediction performance showed the new mixed statistical approach yielded shorter detection times and was more reliable than monitoring future prediction performance only. As a result of this new dual validation mechanism, there were a few false positives and significant delays in identifying any meaningful drift occurrence. A computational complexity analysis confirmed that the new framework will grow linearly with the number of dimensions (i.e., features) being monitored allowing for real-time implementation.

Overall, the results validate the robustness, interpretability, and practical applicability of the proposed framework in dynamic predictive environments.

VII. CONCLUSION

This proposed framework describes a statistically-based system for the detection of real-time concept drift in predictive systems that provide predictions based on varying datasets. The proposed approach combines statistical measures from multiple distributions and continuous monitoring of system performance to create a quantitative relationship between changes in the data and the degradation of the accuracy of the predictive model (concept drift).

The proposed method for concept drift detection (and validation) introduces two sources of validation using two conditions. As such, the use of multiple sources of validation improves the reliability of the method by providing two types of validation and decreasing the number of false positives.

An experimental evaluation of the proposed method, using real-world financial transactions data, found that the proposed method can detect concept drift earlier than alternative performance-based methods, while being computationally efficient.

Statistically rigorous, performance-aware, and modular architectures make the proposed method suitable for use in real-time machine learning applications such as fraud detection, healthcare analytics, and financial risk assessment.

VIII. FUTURE SCOPE

While the proposed framework exhibits outstanding performance, there still exists the potential for future research in several areas.

The existing implementation is currently focused on distribution-based drift detection and could be extended to enable direct detection of conditional distribution shifts, $P(Y|X)$ (true concept drift disrupting decision boundaries).

Also, incorporating adaptive threshold selection methods to dynamically adjust statistical sensitivity values depending on the level of volatility in the environment would enhance the robustness of the framework when dealing with highly-non-stationary environments.

The ability to integrate retraining and adapt at scale through automated pipelines, coupled with reinforcement learning-based adapted models, creates an opportunity for a complete autonomous drift management system to be established.

In addition, validating across a range of domains (healthcare diagnostics, recommendation systems or sensor networks) will provide further understanding of the generalization and practical examples of the scope of application of the findings.

REFERENCES

- [1] J. Gama, P. Medas, G. Castillo, and P. Rodrigues, "Learning with drift detection," in *Advances in Artificial Intelligence – SBIA 2004*, Lecture Notes in Computer Science, vol. 3171. Berlin, Germany: Springer, 2004, pp. 286–295.
- [2] M. Baena-García, J. del Campo-Ávila, R. Fidalgo, A. Bifet, R. Gavalda, and R. Morales-Bueno, "Early drift detection method," in *Proc. 4th Int. Workshop Knowledge Discovery from Data Streams*, 2006, pp. 77–86.
- [3] A. Bifet and R. Gavalda, "Learning from time-changing data with adaptive windowing," in *Proc. 2007 SIAM Int. Conf. Data Mining (SDM)*, Minneapolis, MN, USA, 2007, pp. 443–448.
- [4] F. J. Massey Jr., "The Kolmogorov–Smirnov test for goodness of fit," *Journal of the American Statistical Association*, vol. 46, no. 253, pp. 68–78, Mar. 1951.
- [5] E. M. Lewis, "An introduction to credit scoring," *Fair, Isaac and Company*, San Rafael, CA, USA, Tech. Rep., 1992.
- [6] S. Kullback and R. A. Leibler, "On information and sufficiency," *The Annals of Mathematical Statistics*, vol. 22, no. 1, pp. 79–86, Mar. 1951.
- [7] N. C. Oza and S. Russell, "Online bagging and boosting," in *Proc. 8th Int. Workshop Artificial Intelligence and Statistics (AISTATS)*, 2001, pp. 105–112.
- [8] J. Gama, I. Žliobaitė, A. Bifet, M. Pechenizkiy, and A. Bouchachia, "A survey on concept drift adaptation," *ACM Computing Surveys*, vol. 46, no. 4, pp. 1–37, 2014.

IEEE conference templates contain guidance text for composing and formatting conference papers. Please ensure that all template text is removed from your conference paper prior to submission to the conference. Failure to remove template text from your paper may result in your paper not being published.